


Reasoning tokens per run



Mean latency per run (s)

